

PAXOSLEASE: DISKLESS PAXOS FOR LEASES

MARTON TRENCSENI

CONTENTS

1.	Introduction	1
1.1.	Failure conditions	1
2.	Preliminaries	1
2.1.	Definitions	1
2.2.	Algorithm	2
2.3.	Paxos	3
3.	PaxosLease with perfect clocks	3
3.1.	Acquiring leases	3
3.2.	Extending leases	4
3.3.	Releasing leases	5
3.4.	Learning	6
4.	PaxosLease with skewed clocks	7

1. INTRODUCTION

1.1. **Failure conditions.** PaxosLease handles the following failure conditions:

- (1) Nodes stop and restart
- (2) Network splits
- (3) Message loss
- (4) Message reordering
- (5) In-transit message delays

PaxosLease does not handle Byzantine failures and messages corruption.

2. PRELIMINARIES

2.1. Definitions.

2.1.1. *Actors.* A PaxosLease cell consists of *proposers*, *acceptors* and *learners*. We assume there are n acceptors and any number of proposers and learners. In practice nodes often act as all three, but this is an implementation issue and does not affect the discussion.

2.1.2. *Messages.* Proposers send *prepare request* and *propose request* messages to acceptors, who reply with *prepare response* and *propose response* messages. Proposers send *learn* messages to learners.

The messages have the following structure:

- (1) $\langle \text{prepare request} \rangle = \langle \text{ballot number} \rangle$
- (2) $\langle \text{prepare response} \rangle = \langle \text{ballot number}, \text{accepted proposal} \rangle$
- (3) $\langle \text{propose request} \rangle = \langle \text{proposal} \rangle$
- (4) $\langle \text{propose response} \rangle = \langle \text{ballot number} \rangle$
- (5) $\langle \text{learn} \rangle = \langle \text{value} \rangle$

A *proposal* consists of a *ballot number* and a *value*.

2.1.3. *State.* Acceptors store the following state information:

- (1) *highest ballot promised*: the highest ballot number sent in a prepare response
- (2) *accepted proposal*: the last proposal accepted

2.1.4. *Ballot numbers.* Ballot numbers are positive integers generated by proposers. Different proposers generate different ballot numbers. In implementations, this is achieved by choosing ballots numbers such that *ballot number mod n = proposer id*.

2.1.5. *Chosen value.* A proposal p containing value v is chosen if *accepted proposal* = p for a majority of acceptors.

2.2. Algorithm.

2.2.1. *Basic flow.* A proposer wishing to propose a value first generates a ballot number and sends prepare request messages to a majority of acceptors.

The acceptor, upon receiving a prepare request with *ballot number*, checks whether *ballot number* > *highest ballot promised*. If not, the acceptor does not respond to the message. If yes, the acceptor constructs a prepare response containing the *ballot number* and the currently accepted proposal. If it has not accepted a proposal, it sends special *empty proposal*. The acceptor sets *highest ballot promised* := *ballot number* and then sends the prepare response to the proposer.

The proposer, upon receiving prepare responses from a majority of acceptors, examines the responses. If a majority of nodes responded with empty values, the proposer is free to propose its own value. If not, it must propose the value belonging to the highest *accepted ballot number* among the prepare responses. The proposer sends propose requests containing the generated ballot number and the value to the acceptors who responded.

The acceptor, upon receiving a propose request with proposal p containing *ballot number*, checks its *highest ballot promised*. If *ballot number* < *highest ballot promised* it ignores the message, else it accepts the message by setting *accepted proposal* := p . Finally it sends a propose response with the ballot number.

2.3. Paxos. Augmenting the above rules by two requirements we get Paxos. First, proposers, using stable storage, must guarantee that proposals issued at a later time carry higher ballot numbers. Second, acceptors store their state to stable storage before sending responses.

2.3.1. Consistency. Paxos guarantees that only one value is chosen.

Suppose two proposals p and p' with ballot numbers b and b' and values v and v' are chosen. Since ballot numbers are unique, if $b = b'$ then $p = p'$ and $v = v'$. Suppose $b < b'$. For value v' to be chosen, p' must be proposed.

2.3.2. Lemma. If a proposal p_0 with ballot number b_0 and value v_0 is chosen, then all higher numbered proposals must contain v_0 .

Proof of lemma. Let $p_1 = (b_1, v_1)$ be the proposal with the lowest ballot number such that $b_0 < b_1$. Let A_0 be a majority of acceptors who accepted p_0 and let A_1 be a majority of acceptors who sent prepare responses to p_1 . There must be an acceptor a who is in both sets, who must have *first* accepted p_0 and *later* replied to the prepare request with b_1 . In the prepare reply a indicated that it accepted p_0 , and since p_1 is the lowest numbered proposal after p_0 the proposer could not have received a reply containing a higher numbered proposal, it must have used the value v_0 from proposal p_0 , thus $v_0 = v_1$. Repeating this argument for the next proposal p_2 , if a majority of nodes sent prepare replies, at least one must have contained p_0 or p_1 , thus p_2 must contain the value v_0 . Continuing *ad infinitum* proves the lemma.

3. PAXOSLEASE WITH PERFECT CLOCKS

3.1. Acquiring leases. Augmenting the basic flow with the following rules we get *PaxosLease_{acquire}* (in the perfectly synchronized clock case):

- (1) there is a global maximal lease time T known to all nodes
- (2) upon restart, nodes do not rejoin the network for T time
- (3) for proposals issued within time T , proposals issued at a later time carry higher ballot numbers
- (4) proposers can only propose themselves as lease owners and must honor the global maximal lease time: all values proposed by the k th proposer are of the form "*acquire, k, t_e* " where $t < t_e < t + T$
- (5) messages concerning expired leases are discarded: if an acceptor, proposer or learner receives a message with value "*acquire, k, t_e* " and $t_e < t$ then the message is automatically discarded
- (6) acceptors use a timer and discard their state: before sending a propose response for a value "*acquire, k, t_e* " the node resets its timer to expire at t_e
- (7) if an acceptor's timer expires at t_e , it discards any accepted proposals by setting *accepted proposal* := *empty proposal*

Note that nodes do not have to store their state on stable storage, hence the title "diskless Paxos"!

3.1.1. *Consistency.* $PaxosLease_{acquire}$ guarantees that if value "acquire, k, t_e " is chosen at time t , then no other value is chosen until time t_e .

Suppose two proposals p and p' with ballot numbers b and b' and values $v = \text{"acquire, } k, t_e\text{"}$ and $v' = \text{"acquire, } k', t'_e\text{"}$ are chosen at times t and t' such that $t < t' < t_e < t + T$ and $t < t' < t'_e < t' + T$.

Proposal p was issued within the time interval $(t_e - T, t)$, proposal p' was issued within the time interval $(t'_e - T, t')$. By the above inequalities, the union of these intervals is contained in the interval $(t' - T, t')$, which has length T . Since v' was proposed within T time of t and proposers guarantee not to reuse ballot numbers for T time, if $b = b'$ then $v = v'$.

Let A be a majority of acceptors who accepted p at time t and let A' be a majority of acceptors who accepted p' at time t' . There must be an acceptor a who is in both sets and since $t < t' < t + T$ node a did not restart in the meantime. Acceptors, within time T accept proposals with increasing ballot numbers, thus $b < b'$.

3.1.2. *Lemma.* If a proposal p_0 with ballot number b_0 and value $v_0 = \text{"acquire, } k, t_e\text{"}$ is chosen at time t , then all higher numbered proposals chosen at time t' , such that $t < t' < t_e$, must contain v_0 .

Proof of lemma. See proof of previous lemma and note that a could not have restarted.

3.2. **Extending leases.** $PaxosLease_{acquire, extend}$: The above algorithm can be extended to allow proposers to *extend their leases*. Only the proposer's rule is changed: if a proposer k receives prepare responses from a majority of acceptors and the highest numbered reply contains the value "acquire, k, t_e " such that $t < t_e$ (its lease has not yet expired) then it is free to propose a different value "acquire, k, t'_e " such that $t < t_e < t'_e < t + T$.

3.2.1. *Consistency.* $PaxosLease_{acquire, extend}$ guarantees that if value "acquire, k, t_e " is chosen at time t and value "acquire, k', t'_e " is chosen at time t' such that $t < t' < t_e$ then $k = k'$.

The proof is similar as above. Suppose two proposals p and p' with ballot numbers b and b' and values $v = \text{"acquire, } k, t_e\text{"}$ and $v' = \text{"acquire, } k', t'_e\text{"}$ are chosen at times t and t' such that $t < t' < t_e < t + T$ and $t < t' < t'_e < t' + T$.

Proposal p was issued within the time interval $(t_e - T, t)$, proposal p' was issued within the time interval $(t'_e - T, t')$. By the above inequalities, the union of these intervals is contained in the interval $(t' - T, t')$, which has length T . Since v' was proposed within T time of t and proposers guarantee not to reuse ballot numbers for T time, if $b = b'$ then $v = v'$.

Let A be a majority of acceptors who accepted p at time t and let A' be a majority of acceptors who accepted p' at time t' . There must be an acceptor a who is in both sets and since $t < t' < t + T$ node a did not restart in the meantime. Acceptors, within time T accept proposals with increasing ballot numbers, thus $b < b'$.

3.2.2. Lemma. If a proposal p_0 with ballot number b_0 and value $v_0 = \text{"acquire, } k, t_e\text{"}$ is chosen at time t , then all higher numbered proposals chosen at time t' , such that $t < t' < t_e$, must contain values $\text{"acquire, } k', t'_e\text{"}$ such that $k = k'$.

Proof of lemma. The proof is similar as the previous two lemmas. Let $p_1 = (b_1, v_1)$ be the chosen proposal (at time t' such that $t < t' < t_e$) with the lowest ballot number such that $b_0 < b_1$. Let A_0 be a majority of acceptors who accepted p_0 and let A_1 be a majority of acceptors who sent prepare responses to p_1 . There must be an acceptor a who is in both sets, who must have *first* accepted p_0 and *later* replied to the prepare request with b_1 . In the prepare reply a indicated that it accepted p_0 , and since p_1 is the lowest numbered proposal after p_0 the proposer could not have received a reply containing a higher numbered proposal, it must have used the value v_0 from proposal p_0 , or if the proposer was node k , issued another proposal for itself. Repeating this argument for the next proposal p_2 , if a majority of nodes sent prepare replies, at least one must have contained p_0 or p_1 , thus p_2 must contain a value proposing k as the leaseholder. Continuing *ad infinitum* proves the lemma.

3.3. Releasing leases. In the formulation of PaxosLease (and Paxos) so far, a value was chosen if a majority of acceptors accepted a proposal. In practice, nodes must be able to *learn* that a value was chosen. In classical Paxos, a proposer learns that a value was chosen when it receives *propose responses* for a proposal p from a majority of acceptors. Then it can choose to inform learners by sending them *learn* messages.

In the full formulation of PaxosLease, $PaxosLease_{\text{acquire, extend, release}}$, we change the definition of chosen value: a proposal p containing value v is chosen if the proposer receives propose response messages from a majority of acceptors. Note that this definition is stronger than the previous, so all proofs presented so far remain valid.

Now we can complete PaxosLease by introducing special $\text{"release, } t_r\text{"}$ values. Both the proposer's and acceptor's rules are amended.

- (1) An acceptor, receiving a proposal p with ballot number b (such that *highest ballot promised* $\leq b$) and value $\text{"release, } t_r\text{"}$, if $t_r < t$ must reset its accepted proposal at time t_r : *accepted proposal* $:=$ *empty proposal*.
- (2) Proposer k , at time t learning that its value $\text{"acquire, } k, t_e\text{"}$ was chosen, where $t < t_e < t + T$, may either extend its lease as described above or release its lease at time t' by issuing a proposal with value $\text{"release, } t_r\text{"}$ where $t < t' < t_r < t_e < t + T$.

3.3.1. *Consistency.* $PaxosLease_{acquire,extend,release}$ guarantees that if value "acquire, k, t_e " is chosen at time t and value "acquire, k', t'_e " is chosen at time t' such that $t < t' < t_e$ and proposer k did not issue proposals containing *release* messages in the meantime, then $k = k'$.

The proof is almost the same as above. Suppose two proposals p and p' with ballot numbers b and b' and values $v = \text{"acquire, } k, t_e\text{"}$ and $v' = \text{"acquire, } k', t'_e\text{"}$ are chosen at times t and t' such that $t < t' < t_e < t + T$ and $t < t' < t'_e < t' + T$.

Proposal p was issued within the time interval $(t_e - T, t)$, proposal p' was issued within the time interval $(t'_e - T, t')$. By the above inequalities, the union of these intervals is contained in the interval $(t' - T, t')$, which has length T . Since v' was proposed within T time of t and proposers guarantee not to reuse ballot numbers for T time, if $b = b'$ then $v = v'$.

Let A be a majority of acceptors who accepted p at time t and let A' be a majority of acceptors who accepted p' at time t' . There must be an acceptor a who is in both sets and since $t < t' < t + T$ node a did not restart in the meantime. Acceptors, within time T accept proposals with increasing ballot numbers, thus $b < b'$.

3.3.2. *Lemma.* If a proposal p_0 with ballot number b_0 and value $v_0 = \text{"acquire, } k, t_e\text{"}$ is chosen at time t , then all higher numbered proposals chosen at time t' , such that $t < t' < t_e$, must contain values "acquire, k', t'_e " such that $k = k'$, *provided proposer k did not issue release proposals between t and t' .*

The proof of the lemma is the same as in the $PaxosLease_{acquire,extend}$ case.

3.4. **Learning.** What remains is to inform learners of leases. Only the proposer's algorithm changes:

- (1) If proposer k receives propose response messages from a majority of acceptors containing value $v = \text{"acquire, } k, t_e\text{"}$, then it sends this learn messages containing v to the learners. Learners, upon receiving this value, assume that k holds the lease until time t_e .
- (2) If proposer k , who is currently holding the lease wants to release it at time t_r , it sends learn messages containing the value $v = \text{"release, } t_r\text{"}$ to learners *before issuing the proposal for v* . Learners, upon receiving this value, if they knew k currently holds the value, assume that k holds the lease until time t_r .

If the proposer sends the learn messages containing "release..." values to the learners and then fails, then there may be an interim period (of length less than T time) when, according to the learners, noone holds the lease, but no other proposer will be able to acquire it. This does not affect correctness.

On the other hand, if the proposer does not fail and sends enough propose messages containing "release..." values, other proposers will be able to acquire the lease before the proposer's previously held lease's expire time.

4. PAXOSLEASE WITH SKEWED CLOCKS

So far, we have assumed that all nodes know some global time t . In practice this is not true, so we assume the k th node has a local clock $c_k(t)$ with a maximal skew S such that

- (1) $|c_k(t) - c_j(t)| < S$ for all nodes k, j and
- (2) the global maximal lease time T is chosen such that $S < T$.